

## Final Report

### Findings

The objective of this project was to investigate whether socioeconomic and development indicators could be used to predict national air quality through machine learning techniques (Zhang et al., 2012). To accomplish this objective, data from the Global Air Pollution Dataset were integrated with multiple socioeconomic indicators obtained from the World Bank to create a unified dataset combining environmental measurements with demographic and economic characteristics. Although the original intent of the project was to predict AQI classifications, exploratory analysis demonstrated that modeling continuous AQI values preserved substantially more information and allowed for more meaningful comparisons among regression algorithms. As a result, the project evolved into predicting average AQI values using regression models, providing a richer understanding of the relationships between socioeconomic variables and air quality while remaining aligned with the project's overall research objective.

Exploratory data analysis revealed considerable variation in air quality among countries, as well as meaningful differences in the socioeconomic characteristics associated with those observations. Descriptive statistics indicated substantial variability in AQI values, population density, economic development, greenhouse gas emissions, energy consumption, and industrial activity across the dataset. The visualizations created during the analysis reinforced these findings by demonstrating that simple linear relationships between socioeconomic indicators and AQI were generally weak. For example, scatterplots showed considerable dispersion in the relationship between AQI and industry value added, while the correlation heatmap suggested that many predictor variables exhibited only modest linear associations with air quality. Collectively, these observations indicated that air quality is influenced by multiple interacting factors rather than a single dominant socioeconomic variable, supporting the use of machine learning methods capable of modeling nonlinear relationships.

The visualizations developed during the project also played an important role in guiding subsequent analytical decisions. The correlation heatmap provided an overview of the relationships among the selected variables and confirmed that no single socioeconomic indicator exhibited a strong linear relationship with AQI. The scatterplot comparing AQI and industry value added illustrated that countries with similar levels of industrial activity could experience substantially different air quality outcomes, reinforcing the presence of nonlinear relationships. Finally, the Random Forest feature importance visualization complemented the predictive modeling results by identifying the variables that contributed most to model performance. Together, these visualizations transformed the raw data into meaningful insights and informed the selection of machine learning methods used later in the project.

To evaluate the predictive capability of the engineered dataset, three machine learning models were developed and compared: Multiple Linear Regression, Random

Forest Regression, and Gradient Boosting Regression. Multiple Linear Regression served as a baseline model that estimated linear relationships among the predictor variables, while the two ensemble methods were selected because of their ability to capture complex nonlinear interactions within the data. Model performance was evaluated using the coefficient of determination ( $R^2$ ), Mean Absolute Error (MAE), and Root Mean Squared Error (RMSE), providing complementary measures of predictive accuracy and allowing the strengths and limitations of each modeling approach to be compared objectively.

Among the three models, Random Forest Regression demonstrated the strongest overall performance and was selected as the champion model for the project. It achieved the highest test  $R^2$  while also producing the lowest MAE and RMSE, indicating superior predictive accuracy compared with both Multiple Linear Regression and Gradient Boosting Regression. In addition to its predictive performance, the Random Forest model generated feature importance scores that identified Industry Value Added (% of GDP) as the most influential socioeconomic predictor of national AQI. Other variables, including GDP per capita, energy use per capita, greenhouse gas emissions per capita, urban population percentage, population density, and total population, also contributed to the model, although their relative importance was noticeably lower. The comparative strengths and limitations of each predictive model are summarized in Table 8.1

**Table 8.1**  
*Comparison of Predictive Models*

| Model                        | Strengths                                                                       | Limitations                                                                            | Overall Assessment                                      |
|------------------------------|---------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|---------------------------------------------------------|
| Multiple Linear Regression   | Simple, interpretable baseline model                                            | Unable to capture nonlinear relationships; lowest predictive performance               | Useful benchmark but not appropriate as the final model |
| Random Forest Regression     | Highest $R^2$ , lowest MAE and RMSE; identified influential predictor variables | Less interpretable than linear regression                                              | Selected as the champion model                          |
| Gradient Boosting Regression | Strong predictive capability and handled nonlinear relationships well           | Slightly lower performance than Random Forest; more sensitive to hyperparameter tuning | Competitive alternative but not selected                |

Although Random Forest Regression produced the strongest predictive performance, the overall accuracy of all three models remained moderate rather than exceptionally high. This finding suggests that socioeconomic indicators alone cannot fully explain differences in national air quality. Air pollution is influenced by numerous additional factors that were beyond the scope of the integrated dataset, including

meteorological conditions, topography, seasonal weather patterns, transportation emissions, environmental regulations, and localized pollution sources. Consequently, the project demonstrates that while socioeconomic variables provide meaningful predictive information, they represent only one component of a much larger and more complex environmental system.

Overall, the project successfully demonstrated the value of integrating multiple publicly available datasets and applying machine learning techniques to investigate a complex environmental problem. The analysis identified meaningful relationships between socioeconomic development and air quality while illustrating both the capabilities and limitations of predictive modeling using country-level aggregated national data. Beyond identifying Random Forest Regression as the most effective predictive model, the project established a scalable analytical framework that can be expanded through the incorporation of additional environmental variables, higher-resolution geographic data, and more advanced machine learning techniques. These findings contribute to a better understanding of the factors associated with global air quality and provide a foundation for future research aimed at improving predictive performance and supporting evidence-based environmental decision-making (Gupta et al., 2023).

#### **Review of Ethical Aspects for Your Selected Project**

Throughout the execution of this project, ethical considerations centered primarily on the responsible collection, integration, analysis, and interpretation of publicly available data rather than concerns surrounding personally identifiable information. The datasets used in this research were obtained from reputable public sources, including the Global Air Pollution Dataset available through Kaggle and socioeconomic indicators published by the World Bank. Because all data were aggregated at the national level and contained no individual or sensitive information, the project presented minimal privacy risks. Ethical data science, however, extends beyond privacy alone. Ensuring that data were accurately merged, properly documented, and appropriately attributed was essential for maintaining the integrity of the analysis and allowing the research to be reproduced by others.

Another important ethical consideration involved minimizing bias and avoiding unsupported conclusions. While the machine learning models identified statistically meaningful relationships between socioeconomic variables and average national AQI values, the analysis was carefully interpreted as identifying associations rather than establishing causal relationships. Air quality is influenced by numerous environmental, meteorological, geographic, and policy-related factors that were not incorporated into the final dataset. Presenting the findings without acknowledging these limitations could have resulted in misleading conclusions or overconfidence in the predictive capabilities of the models. By openly discussing the moderate predictive performance and recognizing the limitations of the available data, the project provides a balanced and responsible interpretation of the results.

A further ethical responsibility involved ensuring that the machine learning models were used to support analysis rather than replace human judgment. Although the Random Forest model demonstrated the strongest predictive performance, its results should be viewed as one source of evidence rather than definitive proof of the factors that determine air quality. Predictive models are only as reliable as the data used to train them, and important environmental influences remained outside the scope of this study. Consequently, any policy or planning decisions based on these findings should also incorporate scientific expertise, local environmental knowledge, and additional sources of evidence. Recognizing both the capabilities and limitations of machine learning helps promote responsible use of artificial intelligence in environmental research by emphasizing transparency, accountability, and responsible interpretation of model outputs (Jobin et al., 2019).

Transparency and reproducibility were also central ethical principles throughout the project. Each stage of the workflow, including data acquisition, preprocessing, feature engineering, visualization development, model training, and model evaluation, was documented so that another analyst could replicate the methodology using the same publicly available datasets. Model selection was based on objective performance metrics rather than personal preference, and the Random Forest Regression model was selected because it consistently outperformed the alternative approaches across multiple evaluation measures. Maintaining thorough documentation, clearly explaining methodological decisions, including the transition from AQI classification to continuous AQI prediction, and openly discussing the strengths and limitations of the final models helped ensure that the project adhered to accepted principles of ethical and responsible data science.

### **Review of Success or Completion**

Overall, the project successfully achieved its primary objective of investigating the relationship between socioeconomic indicators and national air quality through the application of machine learning techniques. The planned milestones were completed throughout the course, including defining the research problem, integrating multiple public datasets, conducting exploratory data analysis, developing visualizations, training and evaluating multiple predictive models, and identifying a champion model. The project also satisfied the technical requirements of the capstone by utilizing multiple data sources, performing feature engineering, and comparing several machine learning algorithms using established performance metrics. As a result, the project produced a comprehensive analytical workflow that addressed the original research question while demonstrating the application of data science methodologies learned throughout the program.

Although the project was successful, it also highlighted several opportunities for improvement that became apparent during the analysis. The moderate predictive performance of the models demonstrated that socioeconomic variables alone cannot fully explain variation in national air quality, indicating that additional environmental and meteorological variables would likely improve future models. Similarly, the project evolved from predicting AQI classifications to predicting continuous AQI values after

exploratory analysis revealed that the regression approach would preserve more information and provide greater analytical value. Rather than representing a deviation from the project's goals, this methodological adjustment reflected an evidence-based decision that strengthened the overall quality of the analysis. Taken together, the completed project met its objectives while providing a strong foundation for future research and model refinement. More importantly, it demonstrated how an end-to-end data science workflow can transform publicly available data into meaningful insights that support evidence-based environmental decision-making.

## **Review of KPIs**

The first key performance indicator focused on the successful completion of the project and the development of a complete predictive modeling workflow. This KPI was fully achieved through the completion of each project milestone, beginning with defining the research problem and identifying appropriate datasets, followed by data integration, exploratory data analysis, visualization development, predictive modeling, and final model evaluation. Each stage built upon the previous assignment, resulting in a comprehensive analytical framework that satisfied the technical and academic requirements of the capstone project. Successfully completing this workflow demonstrated the practical application of data analytics concepts and machine learning techniques learned throughout the program.

The second KPI evaluated the predictive performance of the machine learning models developed during the project. Three regression algorithms were trained and compared using Multiple Linear Regression, Random Forest Regression, and Gradient Boosting Regression. Model performance was assessed using  $R^2$ , Mean Absolute Error (MAE), and Root Mean Squared Error (RMSE), allowing each model to be evaluated objectively. Random Forest Regression achieved the strongest overall performance and was selected as the champion model based on its superior evaluation metrics. Although the overall predictive accuracy remained moderate, the comparison of multiple algorithms provided valuable insight into the strengths and limitations of different modeling approaches while demonstrating that nonlinear ensemble methods were better suited for this dataset than traditional linear regression.

The final KPI measured the practical value of the project's findings and their potential contribution to understanding global air quality. This objective was also achieved by identifying socioeconomic variables that exhibited meaningful relationships with average national AQI values and by demonstrating how multiple public datasets can be integrated to support environmental analysis. The feature importance results highlighted industry value added as the strongest predictor among the variables included in the study, while also illustrating that socioeconomic indicators alone cannot fully explain differences in air quality. These findings provide a useful foundation for future research and demonstrate how machine learning can support evidence-based environmental decision-making when combined with additional environmental, meteorological, and geographic data.

## **Review of Lessons Learned**

One of the most significant lessons learned during this project was the importance of data preparation and feature engineering. While developing the predictive models was ultimately the primary objective, a substantial portion of the project involved locating compatible datasets, cleaning and standardizing the data, resolving differences in country names and formats, and engineering meaningful socioeconomic features. This experience reinforced that the quality of a machine learning model depends heavily on the quality of the data used to train it. Investing time in careful preprocessing and validation created a stronger foundation for the subsequent analysis and improved confidence in the results.

A second lesson learned was that machine learning projects often evolve as new insights emerge during the analytical process. The project initially focused on predicting AQI classifications, but exploratory analysis revealed that predicting continuous AQI values would preserve more information and allow for a more meaningful comparison of regression algorithms. Rather than adhering rigidly to the original plan, the project adapted its methodology based on evidence gathered during the analysis. This experience demonstrated that successful data science requires flexibility and a willingness to revise an approach when the data indicate that an alternative methodology is more appropriate.

The final lesson learned was that predictive performance alone does not determine the value of a machine learning project. Despite Random Forest Regression having the strongest results among the evaluated models, the overall predictive accuracy remained moderate, illustrating that no model can fully compensate for missing or incomplete information. Rather than viewing this as a failure, the project highlighted the importance of understanding the limitations of available data and interpreting model results responsibly. Identifying those limitations, explaining their impact on model performance, and recognizing opportunities for future improvement are essential components of conducting meaningful and trustworthy data science research.

## **Recommendations for Future Analysis**

Future research should focus on expanding the dataset to include additional environmental and meteorological variables that have well-established relationships with air quality. While this project successfully demonstrated that socioeconomic indicators contribute to explaining differences in national AQI, these variables represent only one aspect of a much more complex system. Incorporating features such as temperature, humidity, precipitation, wind speed, atmospheric pressure, wildfire activity, traffic density, and satellite-derived pollution measurements could substantially improve predictive performance by capturing the environmental conditions that directly influence pollutant formation and dispersion. Additionally, collecting data at regional or city-level resolutions rather than national averages would allow future models to account for local variability that is masked through country-level aggregation. Previous research has similarly demonstrated that meteorological conditions play a significant role in shaping

air pollution patterns and should be considered when developing predictive models of air quality (Zhang et al., 2012).

Future analyses should also explore additional machine learning techniques and more comprehensive model development strategies. Although Random Forest Regression produced the strongest performance among the models evaluated in this project, algorithms such as XGBoost, LightGBM, and deep neural networks may better capture complex nonlinear relationships within larger and more detailed datasets. Future studies could also investigate automated hyperparameter optimization, ensemble model stacking, and cross-validation strategies to further improve predictive performance. Expanding the analysis to include time series forecasting would allow researchers to examine how air quality changes over time and evaluate the effects of seasonal patterns, economic development, or policy changes on long-term environmental conditions.

Finally, future research should focus on translating predictive modeling results into practical tools that support environmental decision-making. The analytical framework developed during this project demonstrates that integrating multiple public datasets can provide meaningful insight into factors associated with air quality, but the greatest value lies in applying those insights to real-world problems. Interactive dashboards, geospatial visualizations, and decision support systems could enable policymakers and environmental agencies to monitor changing conditions, identify areas of concern, and evaluate the potential impact of environmental policies. Additionally, future work could revisit AQI classification alongside continuous AQI prediction using richer environmental datasets, allowing researchers to compare regression and classification approaches within a more comprehensive modeling framework. Building upon the foundation established in this project would not only improve predictive accuracy but also strengthen the ability of machine learning to support evidence-based environmental policy, public health planning, and long-term sustainability efforts.

## References

- Breiman, L. (2001). *Random forests*. *Machine Learning*, 45(1), 5–32. <https://doi.org/10.1023/A:1010933404324>
- Byrne, G. (2007). *A statistical primer: Understanding descriptive and inferential statistics*. *Evidence Based Library and Information Practice*, 2(1), 32–47. <https://doi.org/10.18438/B8FW2H>
- Cleveland, W. S., & McGill, R. (1984). *Graphical perception: Theory, experimentation, and application to the development of graphical methods*. *Journal of the American Statistical Association*, 79(387), 531–554. <https://doi.org/10.1080/01621459.1984.10478080>
- Fernandez, N. F., Gundersen, G. W., Rahman, A., Grimes, M. L., Rikova, K., Hornbeck, P., & Ma'ayan, A. (2017). *Clustergrammer, a web-based heatmap visualization and analysis tool for high-dimensional biological data*. *Scientific Data*, 4(1), 170151. <https://doi.org/10.1038/sdata.2017.151>
- Friedman, J. H. (2001). *Greedy function approximation: A gradient boosting machine*. *The Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>
- Gupta, N. S., Mohta, Y., Heda, K., Armaan, R., Valarmathi, B., & Arulkumaran, G. (2023). *Prediction of air quality index using machine learning techniques: A comparative analysis*. *Journal of Environmental and Public Health*, 2023 <https://doi.org/10.1155/2023/4916267>
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An introduction to statistical learning: With applications in R (2nd ed.)*. Springer. <https://doi.org/10.1007/978-1-0716-1418-1>
- Jobin, A., Ienca, M., & Vayena, E. (2019). *The global landscape of AI ethics guidelines*. *Nature Machine Intelligence*, 1(9), 389–399. <https://doi.org/10.1038/s42256-019-0088-2>
- Kaneko, H. (2023). *Interpretation of machine learning models for data sets with many features using feature importance*. *ACS Omega*, 8(25), 23218–23225. <https://doi.org/10.1021/acsomega.3c03722>
- Muzdadid, H. A. (2023). *Global air pollution dataset* [Data set]. Kaggle. <https://www.kaggle.com/datasets/hasibalmuzdadid/global-air-pollution-dataset/data>
- Rybarczyk, Y., & Zalakeviciute, R. (2018). *Machine learning approaches for outdoor air quality modelling: A systematic review*. *Applied Sciences*, 8(12), 2570. <https://doi.org/10.3390/app8122570>
- World Bank Group. (2023). *World development indicators: Population, total (SP.POP.TOTL)*. <https://databank.worldbank.org/source/world-development-indicators>
- Zhang, Y., Bocquet, M., Mallet, V., Seigneur, C., & Baklanov, A. (2012). *Real-time air quality forecasting, part I: History, techniques, and current status*. *Atmospheric Environment*, 60, 632–655. <https://doi.org/10.1016/j.atmosenv.2012.06.031>